

Divide-and-conquer and randomization

1 Tiling by L-shapes

We consider the problem of tiling a *punctured* grid consisting of $2^n \times 2^n$ squares via what we call *L-shapes*. *Punctured* means that one square of the grid, called the *hole*, is missing. An *L-shape* is any set of 3 unit squares that are inside a (necessarily unique) 2×2 square. The integer n is called the *size* of the grid. The file `L_tiling.py` on Moodle contains the skeleton code that you need to complete, and the file `test_L_tiling.py` contains the code to test your solutions.

Each square of a grid is identified by consecutive cartesian coordinates, with x -coordinates increasing from left to right and y -coordinates increasing from bottom to top. The left image in Figure 1 shows an example of a punctured grid of size 2. We call the bottom left corner of a grid its *origin*, which does not necessarily need to be $(0, 0)$. In fact, it is convenient to allow the origin to be any point with non-negative coordinates. That is, for two integer $i, j \geq 0$, the x -coordinates of the grid at origin (i, j) of size n range from i to $i + 2^n - 1$, and the y -coordinates range from j to $j + 2^n - 1$.

Given integers n, i, j, a, b , we define the *punctured grid of type* (n, i, j, a, b) to be the $2^n \times 2^n$ grid with origin (i, j) and with the hole at (a, b) .




Figure 1: Left: the punctured 4×4 grid G with the hole at position $(1, 3)$. Right: an *L*-tiling of G .

We are interested in tiling a punctured grid, that is, cover it with non-overlapping *L*-shapes. The right image of Figure 1 shows an example of such a tiling. In Python, we represent a tiling as a list of triples (the *L*-shapes) of integer pairs. In Figure 1, the list is (up to reordering) $[(1, 1), (2, 2), (2, 1)], [(0, 0), (0, 1), (1, 0)], [(3, 1), (3, 0), (2, 0)], [(0, 2), (0, 3), (1, 2)], [(3, 2), (2, 3), (3, 3)]$.

In order to compute a valid *L*-tiling of a punctured grid G , we use the following divide-and-conquer strategy, also sketched in Figure 2. If the size n of G is zero, return. If n is greater than zero, we decompose G into 4 quadrants (each a grid of size $n - 1$). We call the *marked quadrant* the one containing the hole of G . Let the *middle* M of G be the *L*-shape obtained from the central 2×2 square X by removing the unique square of X that belongs to the marked quadrant (left and central image of Figure 2). We add M to our current tiling. Then we puncture the marked quadrant to have the same hole as G , whereas the other 3 quadrants are punctured at their unique unit square belonging to X (right image of Figure 2). Last, we apply the tiling procedure to each of the 4 quadrants recursively.

Question 1. Complete the function `middleL(n, i, j, a, b)`, which, given a punctured grid of type (n, i, j, a, b) with size $n \geq 1$, returns the middle of the grid as a list $[(x1, y1), (x2, y2), (x3, y3)]$. (Any order of the list is fine.)




Figure 2: The divide-and-conquer strategy to compute an L -tiling of a punctured $2^n \times 2^n$ grid.

Question 2. Complete the function `lower_left_hole(n, i, j, a, b)`, which returns the coordinates `k, l` of the hole in the lower left quadrant of the punctured grid of type (n, i, j, a, b) . Complete analogously the functions for the other three quadrants.

Question 3. Complete the recursive function `tile(n, i, j, a, b)`, which contains a list L -shapes that form a valid tiling of the punctured grid of type (n, i, j, a, b) .

Once this function passes the test, you can uncomment the function `display_tiling_with_random_hole(n)` in `test_L_tiling.py`, which displays the L -tiling of the $2^n \times 2^n$ grid punctured at a square chosen uniformly at random.

2 Randomized MinCut

We now implement a randomized algorithm to compute the minimum cut (mincut) of a multigraph. A multigraph G is a graph where there can be several edges between any pair of vertices. There are no self-loops in a multigraph, that is, edge connects two *distinct* vertices. Figure 3 shows an example (for instance vertices c and g are connected by two edges). The *degree* of a vertex is the number of its incident edges, taking the multiplicity of edges into account. For instance, in Figure 3, the vertex c has degree 4 (and has 3 neighbors). Further, the graph has a mincut of size 3, obtained by taking the partition $V = \{a, c, d, g\} \cup \{b, e, f, h\}$.




Figure 3: A multigraph and its unique mincut.

The file `MultiGraph.py` on Moodle contains the skeleton code that you need to complete, and the file `test_mincut.py` contains the code to test your solutions. The file `MultiGraph.py` contains the class `MultiGraph`, whose instances represent a multigraph. Each instance `M` has an attribute `adj`, which is a dictionary that maps each vertex v its neighbors, and each such neighbor w is mapped to the number of edges between v and w . More precisely:

- Testing if `x` is a vertex of `M` is done by `if x in M.adj`. If the test occurs in a method of the class `MultiGraph`, then this test may need to be written as `if x in self.adj`.

- If x and y are two vertices, testing if x and y are connected by at least one edge is done by `if y in M.adj[x]`. If true, `M.adj[x][y]` returns the number of edges connecting x and y .
- Enumerating the set of pairs (neighbor, number of edges to that neighbor) for a vertex x can be done with `for (y,n) in M.adj[x].items()`.

Further, there is an attribute `deg` such that if x is a vertex of M , then `deg[x]` returns the degree of x . The constructor of `MultiGraph` receives a list, such as `L = [3, [['a', 'b', 1], ['c', 'b', 2], ['a', 'c', 4]]]`, where `L[0]` returns the number of vertices of the multigraph to create, and `L[1]` returns the list of edges, including multiplicities (in this example, there are 4 edges connecting `a` and `c`). An example of such a list is given by `L_tutorial` at the beginning of `test_mincut.py`, which represents the multigraph of Figure 3. Last, each instance of `MultiGraph` has a method `display`, which outputs the graph in text form to the console.

For the convenience of testing, the function `two_clique_graph(n)` in `test_mincut.py` returns (written in this list format) a graph on $2n$ vertices, for which we know the unique mincut. This graph consists of two connected cliques (that is graphs that are fully connected). More precisely, there is a clique for vertices with labels in $[1..n]$, and there is another clique for vertices with labels in $[n+1..2n]$. In addition, there are $n-2$ edges, which, for i from 1 to $n-2$, connect vertex i to vertex $2n-i+1$ (see Figure 4 for an illustration). This graph has mincut size $n-2$, which is uniquely obtained for the partition $[1..n] \cup [n+1..2n]$.




Figure 4: The two-clique graph for $n = 5$.

We now focus on the contraction algorithm seen in class. Recall that the operation of *contracting an edge* $\{u, v\}$ in a multigraph M is defined as follows: The *contracted multigraph* $M/\{u, v\}$ is obtained from M by merging u with v , keeping u as the name of the merged vertex, deleting all the edges connecting u to v (that have become self-loops), and redirecting all edges from v to u . We say that u has *absorbed* v . See Figure 5 for examples.




Figure 5: A multigraph M and the two contracted multigraphs $M/\{e, b\}$ and $M/\{c, d\}$.

Question 4. Complete the method `contract(self, i, j)` in the class `MultiGraph`, which contracts an edge $\{i, j\}$. Do not forget to delete the self-loops arising from the contraction and to update the attributes `deg` and `adj`. Deleting edges between two vertices i and j of a multigraph M can be done with `del M.adj[i][j]` and `del M.adj[j][i]`. Using `del`, it is also possible to delete the entire entry of a node in the dictionary `adj` or `deg`.

A second ingredient in the randomized mincut algorithm is to be able to draw an edge uniformly

at random in a multigraph. To this end, we write a general-purpose function that draws a random element in a set of weighted elements.

Question 5. Complete the method `random_element(dict)`, which is given a dictionary `dict` whose values are positive integers, denoting the weights of each element, and returns a random element with probability proportional to the weight of the element. For instance, for `dict = {'a': 2, 'b': 1, 'c': 4}`, a call to `random_element(dict)` should return `'a'` with probability $2/7$, return `'b'` with probability $1/7$, and return `'c'` with probability $4/7$.

Question 6. Complete the method `random_vertex(self)`, which returns a random vertex of `self`, where the weight of each vertex is its degree. For instance, in Figure 3, the sum of all vertex degrees is 34 (twice the number of edges), and the vertex `'b'` has degree 5. Hence, a call to `self.random_vertex()` should return `'b'` with probability $5/34$.

Further, complete the method `random_edge(self)`, which returns an edge `(i, j)` taken uniformly at random (think of a random edge (i, j) as obtained from a random vertex i and then a random neighbor j of i). For instance, in the graph of Figure 3, there are 17 edges in total, two edges between `'c'` and `'g'`, and one edge between `'a'` and `'d'`. Hence, for a call to `self.random_edge()`, the probability that the output is `('c', 'g')` or `('g', 'c')` should be $2/17$, and the probability that the output is `('a', 'd')` or `('d', 'a')` should be $1/17$.

For a multigraph M with n vertices, a random cut of M is obtained as follows: For i from 1 to $n - 2$, we choose an edge uniformly at random in the current multigraph (which has $n - i + 1$ vertices) and contract it. At the end, there remain two vertices a, b and a bunch of $k \geq 1$ edges connecting them. This naturally corresponds to a partition (S, \bar{S}) of cutsizes k (note that S consists of a and all the vertices that have been absorbed by a).

Question 7. Complete the method `random_cut(m)`, which returns a random cut of the multigraph `m` after contracting $n - 2$ random edges as described above. The format of the output should be `[c, L]`, where `c` is the cutsize and `L` is the list of vertices in one of the two parts of the partition (S, \bar{S}) of the cut. In order to return `L`, you should maintain at each step a dictionary `partition` such that for each `x` not already absorbed, `partition[x]` is the list of vertices that have been already absorbed by `x`. At the end, `partition` has just two keys, `a, b`, corresponding to the two remaining vertices. Thus, and one can return `L` as `[a] + partition[a]`.

The probability that `random_cut(m)` returns a mincut of a multigraph `m` with n vertices is at least $\frac{2}{n(n-1)}$. As seen in class, we can boost the success probability by repeating the process k times, and output a cut of smallest size among the k experiments. Then, the probability of returning a mincut is at least $1 - (1 - \frac{2}{n(n-1)})^k$. Hence for any fixed $p \in (0, 1)$, this probability is guaranteed to be at least $1 - p$ for $k = \lceil \ln(p) / \ln(1 - 2/(n(n-1))) \rceil$ (which is $O(n^2 \log(1/p))$).

Question 8. Complete the function `mincut_karger(L, p)`, which receives as input a multigraph `L` (in the list format) and a floating-point number `p` $\in (0, 1)$ and returns a cut of the multigraph such that the returned cut is a mincut with probability at least $1 - p$. In order to test your code, you can run `test_mincut_karger()`.

Fun fact. You can run the method `test_mincut_brute()` if you want to see how slowly a brute-force solution is. This method takes an optional parameter, which determines the value n of the graph from Figure 4. By default, the method chooses $n = 9$.